

# Chipyard Fundamentals

Jerry Zhao

— eley

jzh@berkeley



Berkeley  
Architecture  
Research

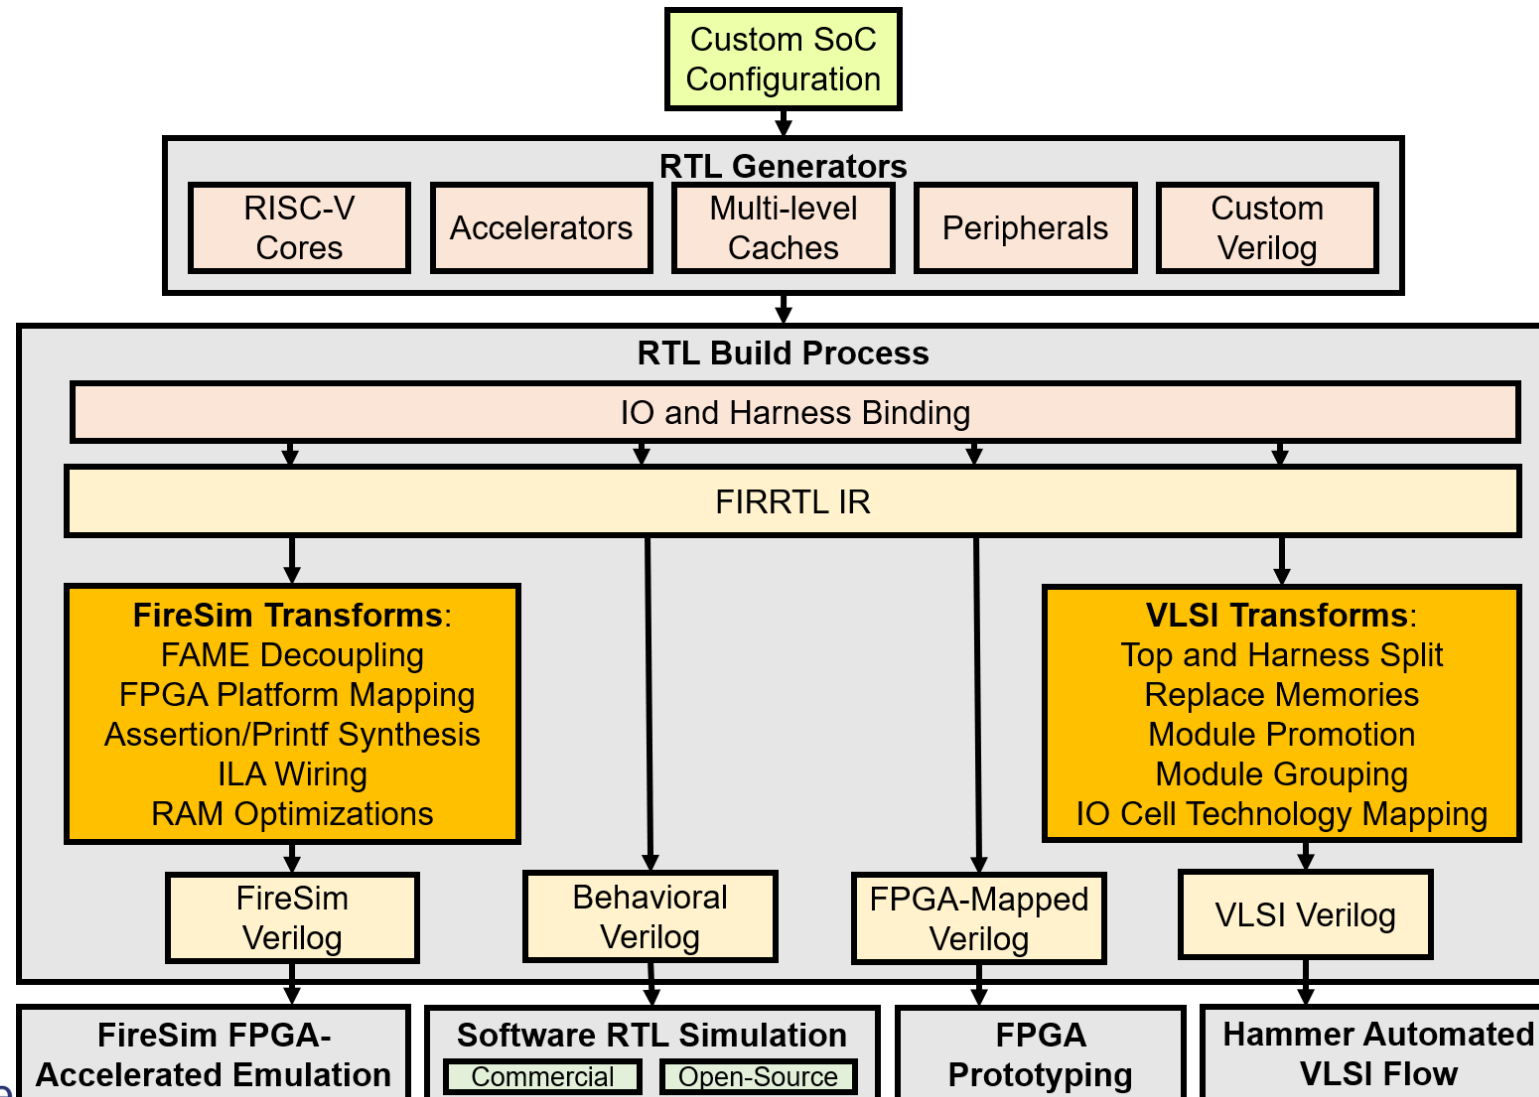
# Outline




- Chipyard Framework Architecture + Integration
- ML Tile
- GP Tile
- Memory subsystem



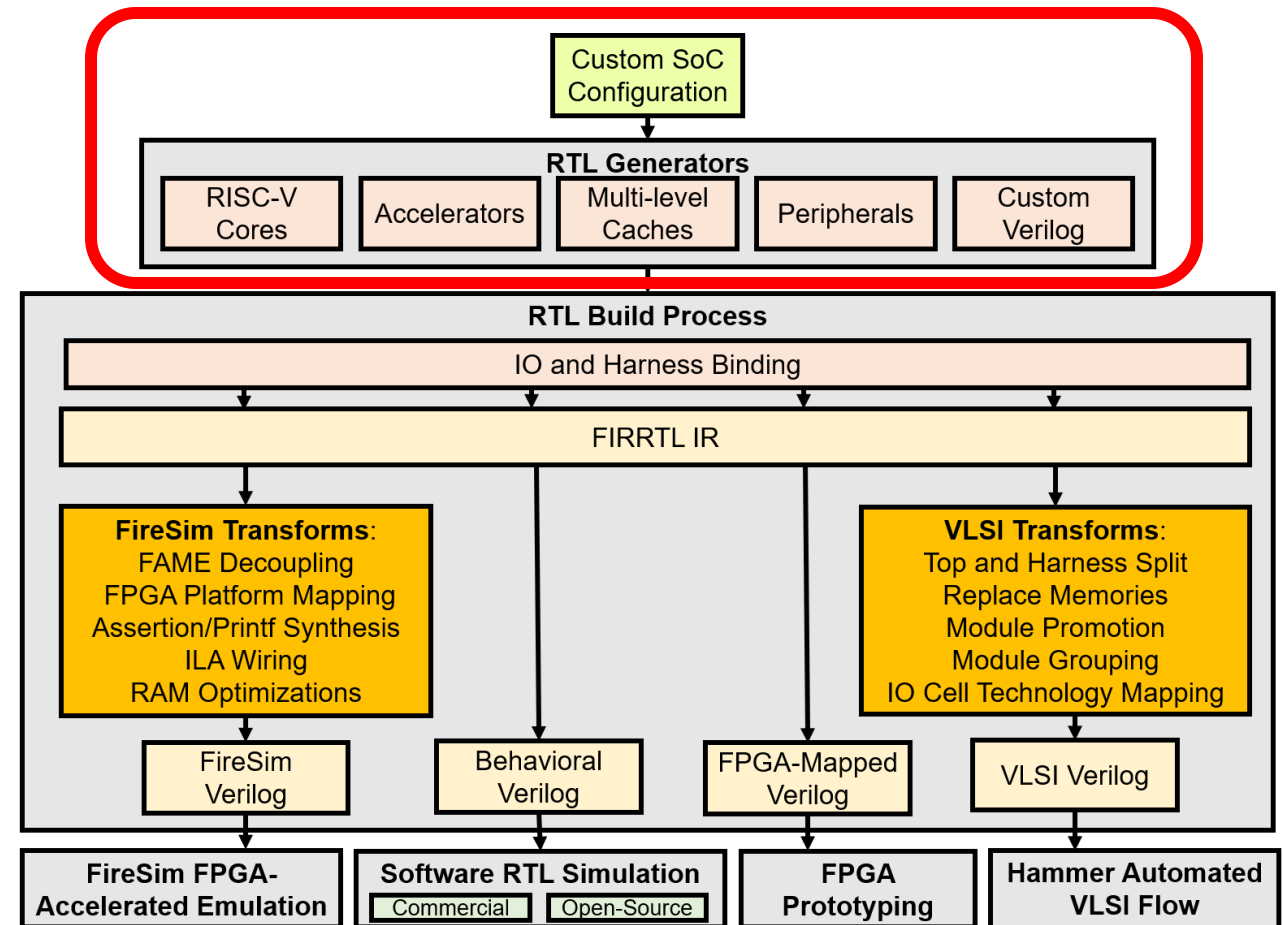
# How is this integrated? Generators!



# How is this integrated? Generators!



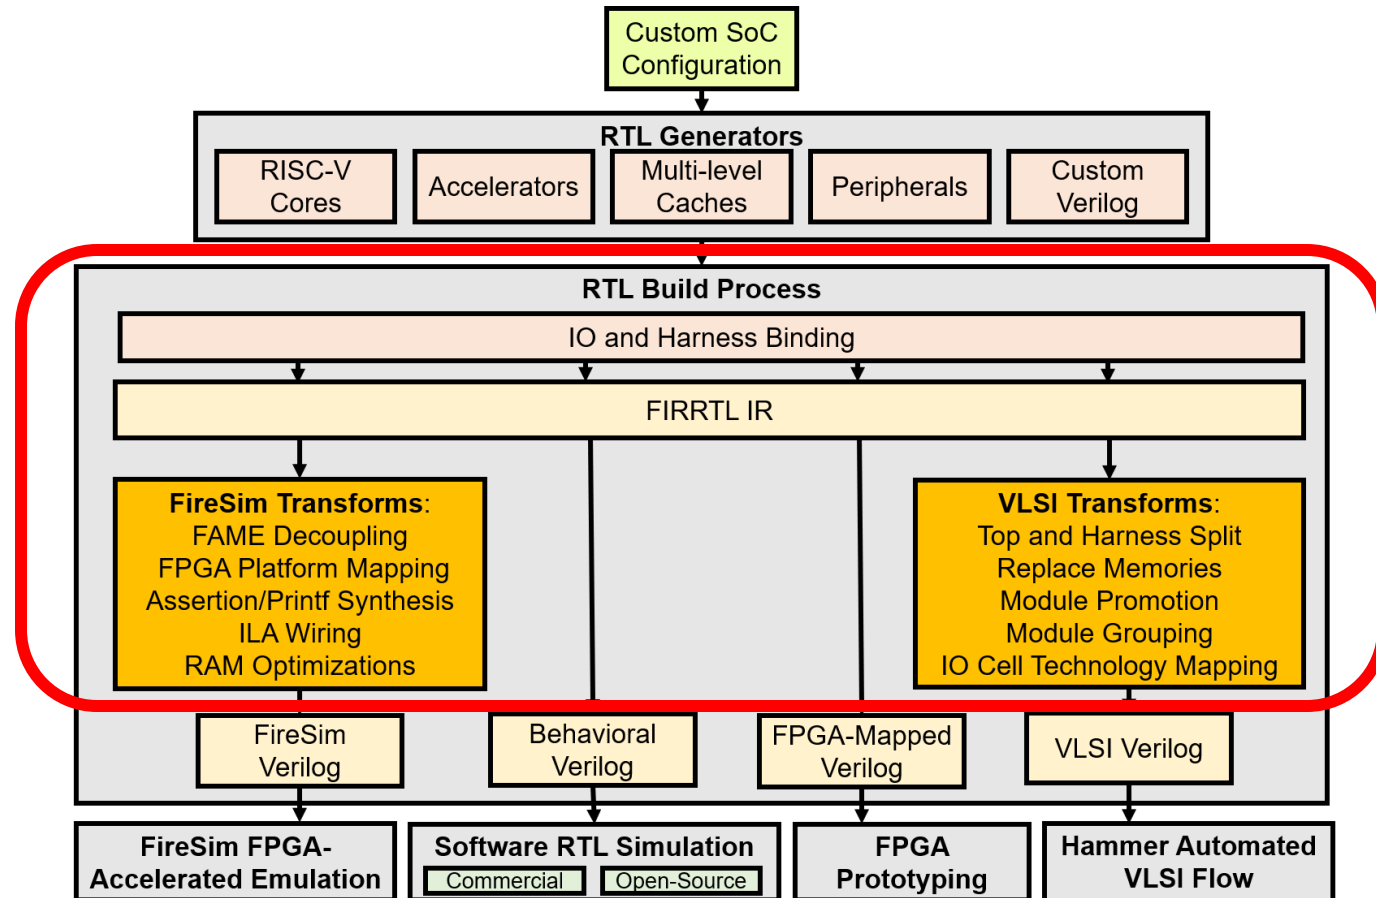
- Everything starts from a generator configuration
- Generators written in Chisel
- Generators can integrate third-party Verilog instance IP



# How is this integrated? Generators!



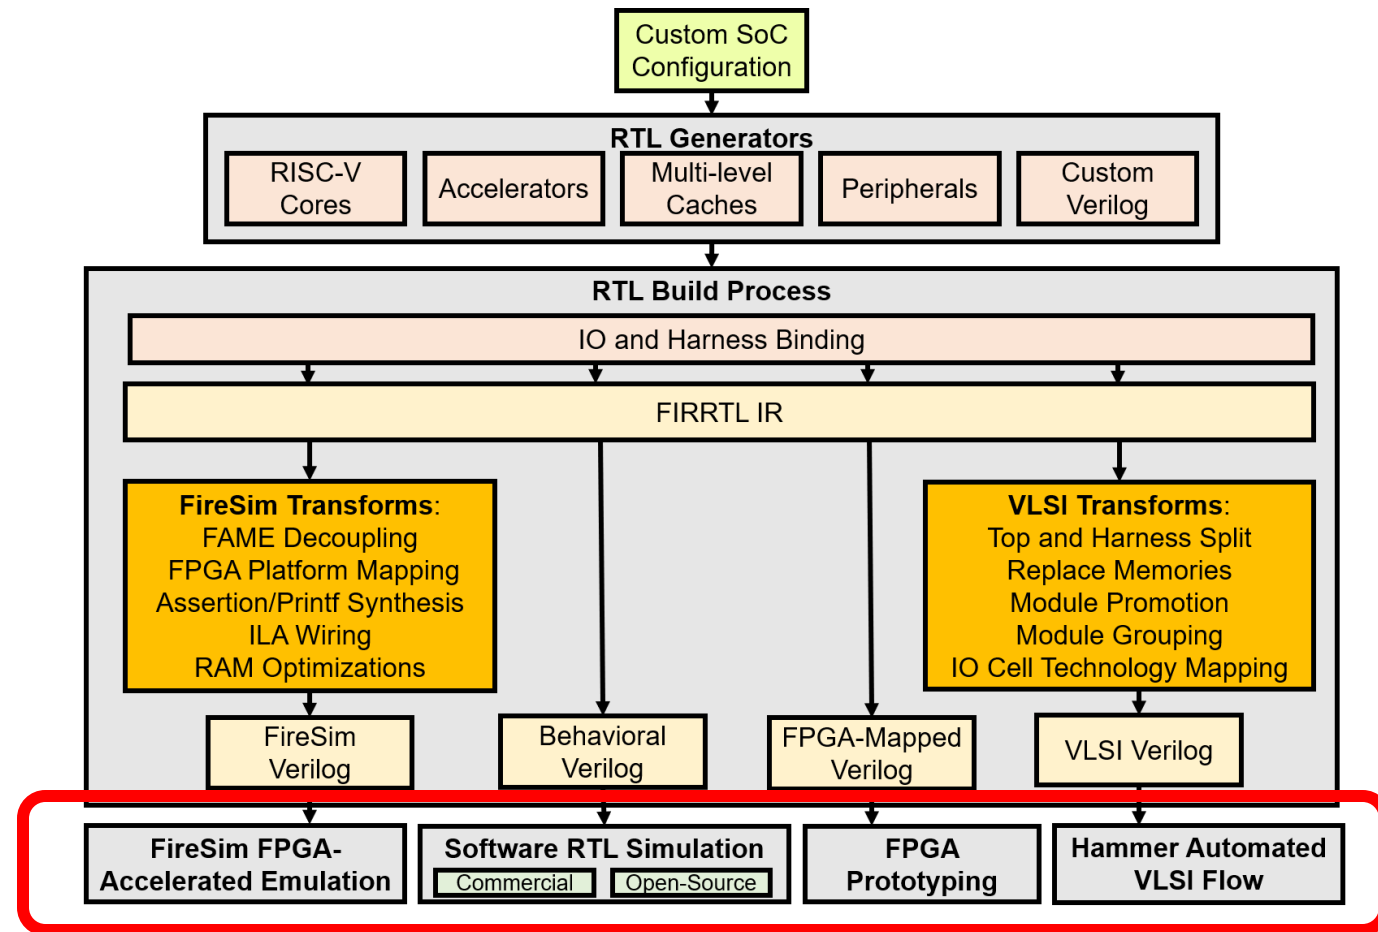
- Elaboration and Transformation
- Internals: FIRRTL – IR enables automated manipulation of the hardware description
- Externals: I/O and Harness Binders – pluggable interface functions enable automated targeting of different external interface requirements



# How is this integrated? Generators!



- Design flows
  - Software RTL Simulation
  - FPGA-Accelerated Emulation
  - FPGA Prototyping
  - VLSI Fabrication



# Chipyard Learning Curve



## Advanced-level


- Configure custom IO/clocking setups
- Develop custom FireSim extensions
- Integrate and tape-out a complete SoC

## Evaluation-level

- Integrate or develop custom hardware IP into Chipyard
- Run FireSim FPGA-accelerated simulations
- Push a design through the Hammer VLSI flow
- Build your own system

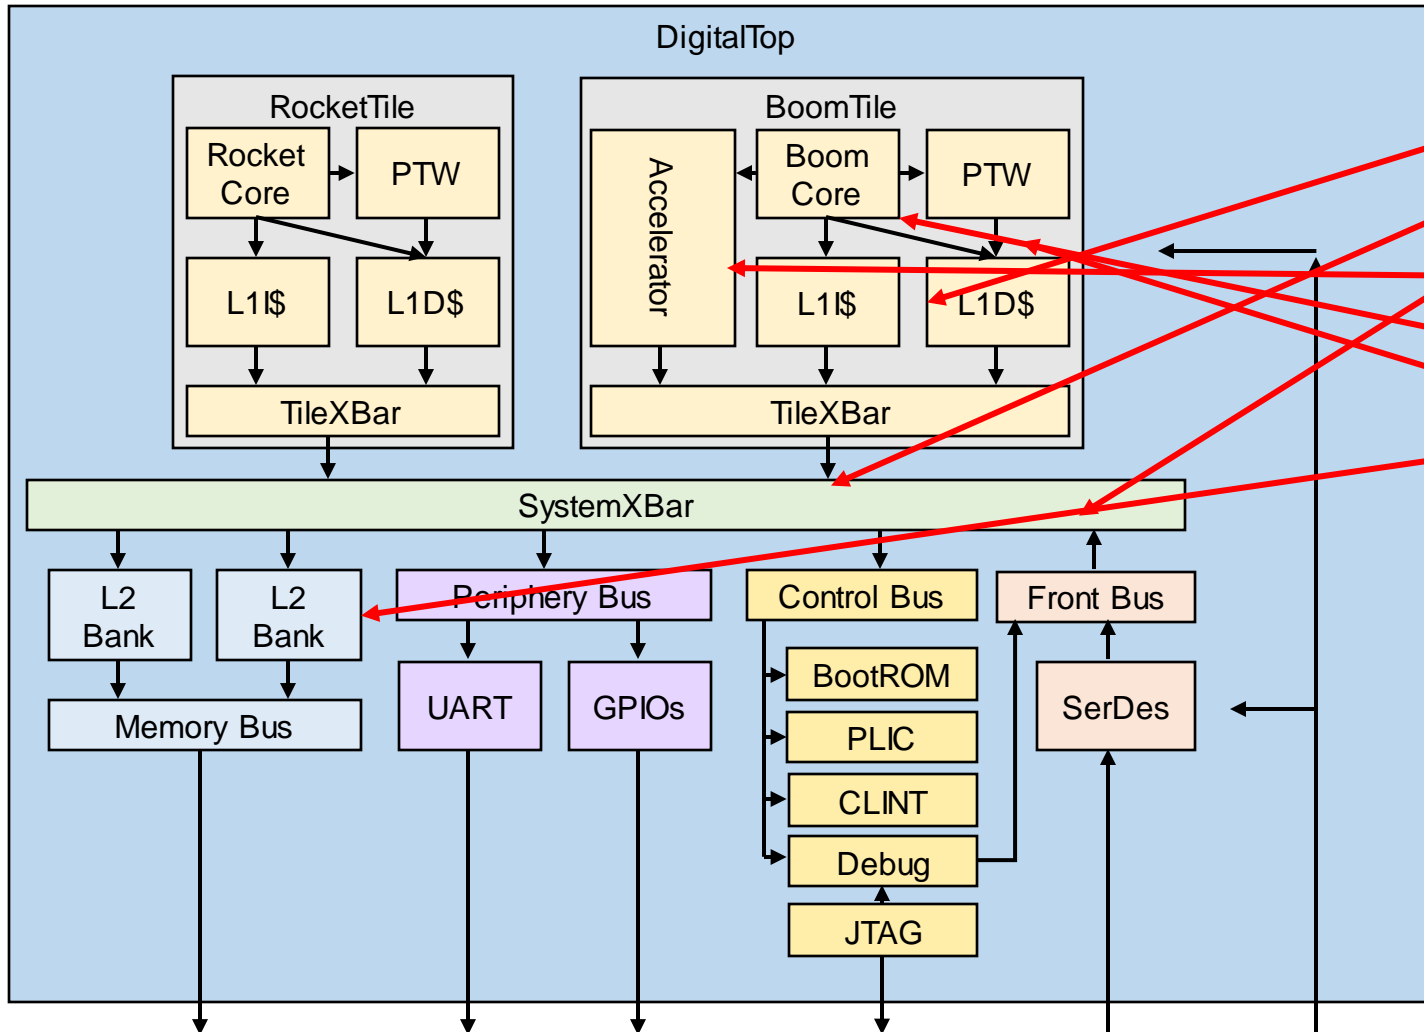
## Exploratory-level

- Configure a custom SoC from pre-existing components
- Generate RTL, and simulate it in RTL level simulation
- Evaluate existing RISC-V designs





# Highly Parameterized Configurations



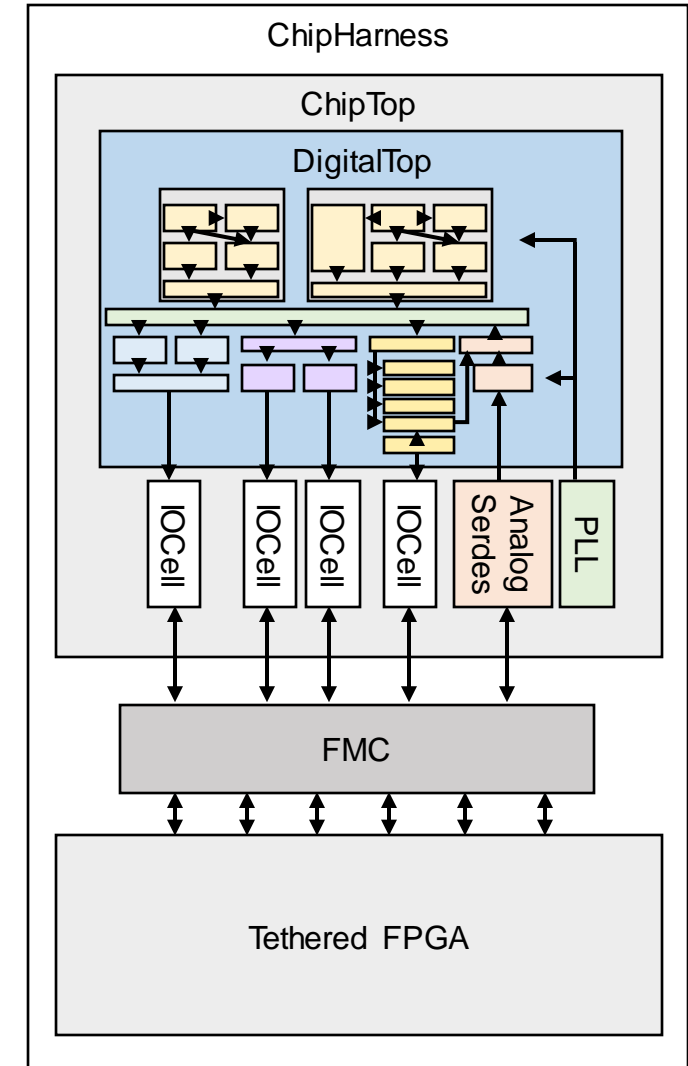
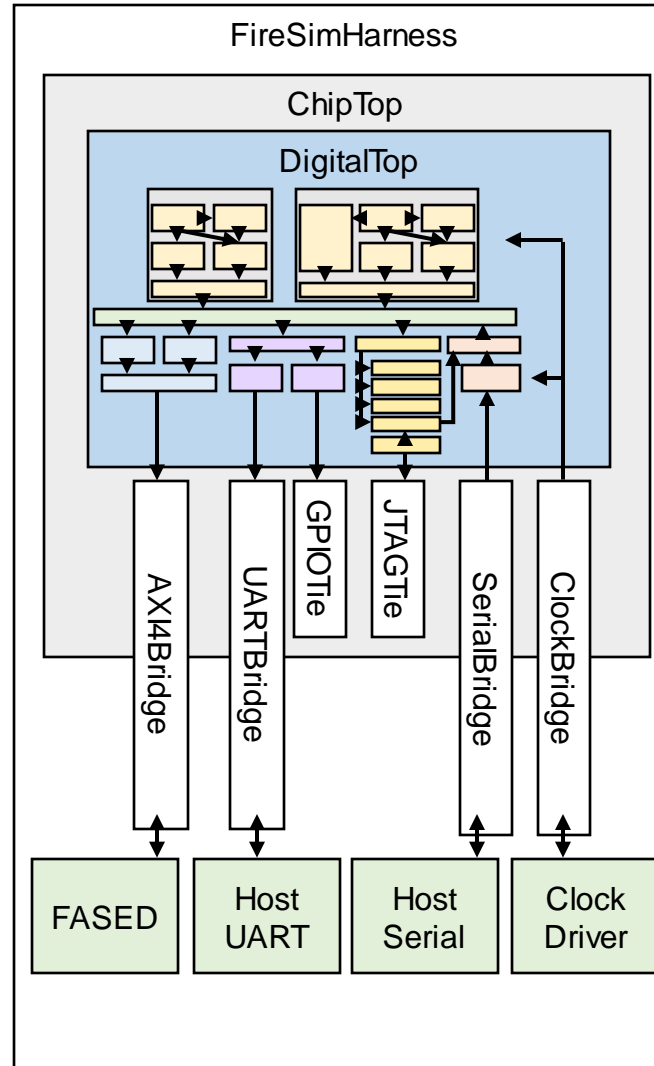
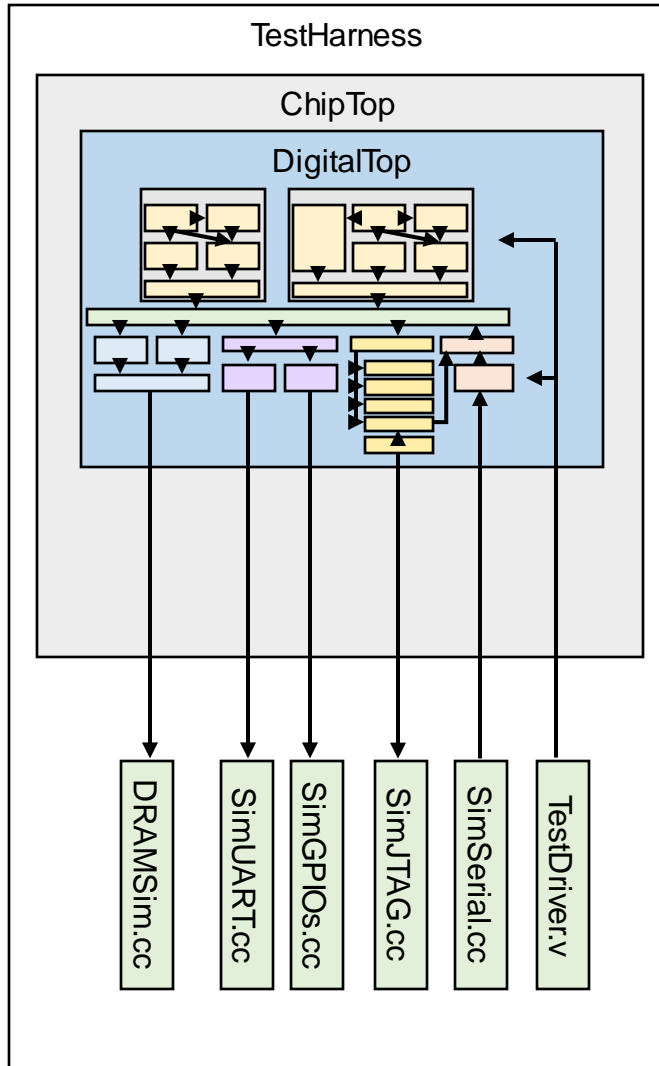
```
class CustomConfig extends Config(  
  new WithL1CacheWays(4) ++  
  new WithAsyncTiles ++  
  new WithRingSystemBus +  
  new WithFPGemmini ++  
  new With3WideBooms ++  
  new WithL2TLBs(512) ++  
  new WithL2Banks(4) ++
```

```
  new WithDefaultGemmini ++  
  new WithNRocketCores(1) ++  
  new WithNBoomCores(1) ++  
  new WithBootROM ++  
  new WithUART ++  
  new WithJtagDTM ++  
  new WithGPIOs ++  
  new WithInclusiveCache(512) ++
```

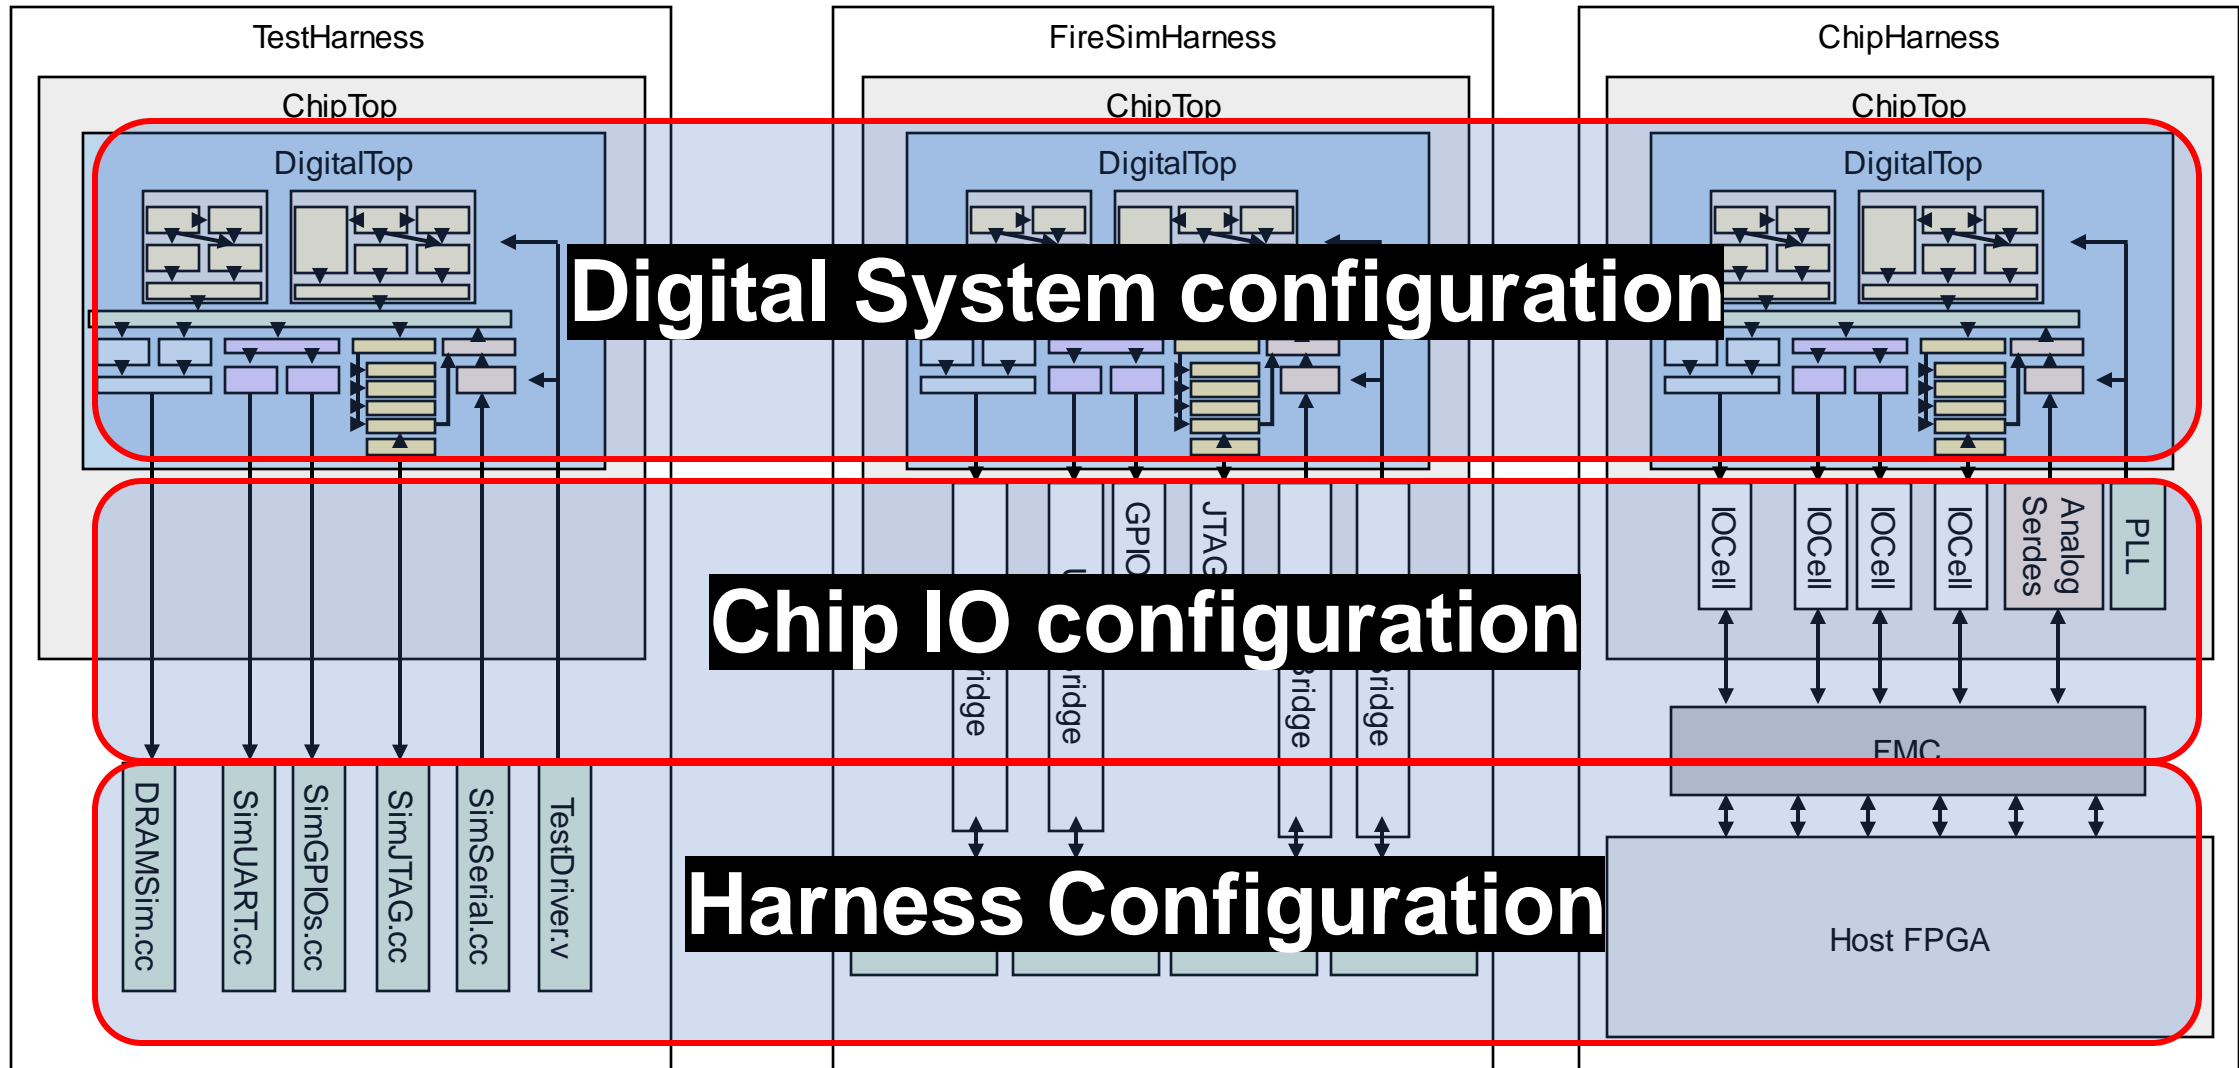




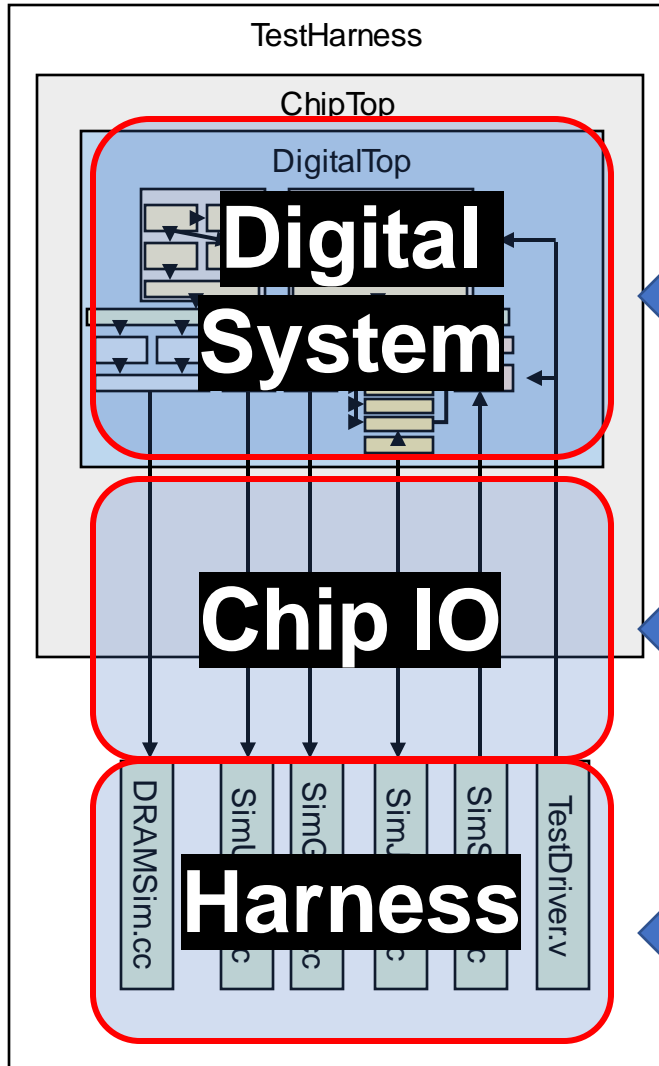
# Multipurpose



# Multipurpose



# A Complete Config



```
class CustomConfig extends Config(  
  new WithDefaultGemmini ++  
  new WithNRocketCores(1) ++  
  new WithNBoomCores(1) ++  
  new WithBootROM ++  
  new WithUART ++  
  new WithJtagDTM ++  
  new WithGPIOs ++  
  new WithInclusiveCache(512) ++  
  
  new WithPassThroughIOs ++  
  
  new WithDRAMSim ++  
  new WithSimUART ++  
  new WithSimJTAG ++  
  new WithSimSerial
```

)





# Using Chipyard Diplomatic Clocking

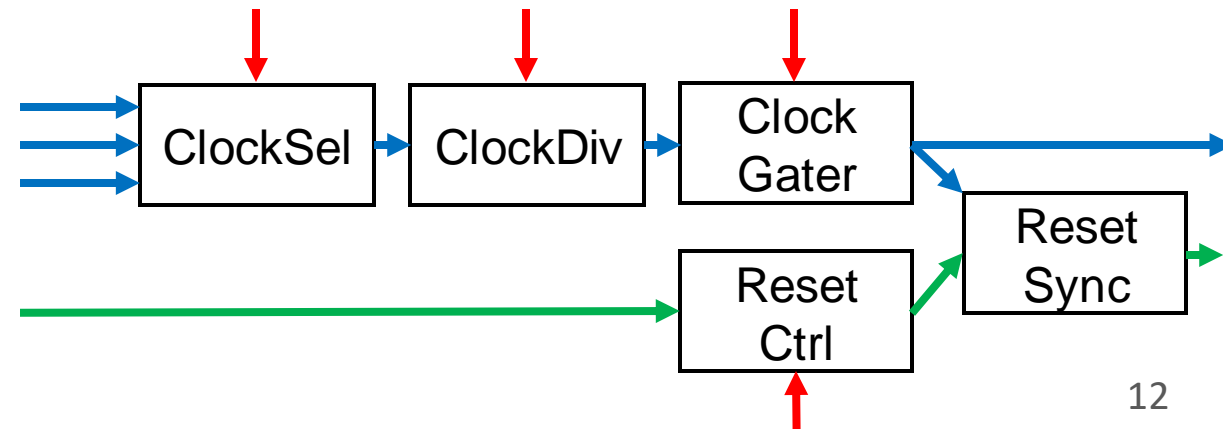
Clock and reset distribution is hard

- Many clock sources
- Programmable clock/reset control
- Dangerous source of bugs

## Chipyard Diplomatic Clocking

- Leverage RocketChip diplomatic clock graphs
- Pull out clock graph out of digital core logic
- Describe clock/reset distribution as graph of nodes

```
(aggregator
  := ClockGroupFrequencySpecifier(p(ClockFrequencyAssigners
  := ClockGroupCombiner("uncore", "implicit", "sbus", "pbu
  := ClockGroupResetSynchronizer()
  := clockTap.clockNode
  := TileClockGater(chiptop.lazySystem.asInstanceOf[BaseSub
  := resetSetterResetProvider
  := fbusClockDiv.clockNode
  := ClockGroupCombiner("uncore", "uncore", "fbus")
  := testchipClockDiv.clockNode
  := testchipClockSel.clockNode)
slowClockSources.foreach(testchipClockSel.clockNode := _)
diffClockSources.foreach(testchipClockSel.clockNode := _)
pllClockSources.foreach(testchipClockSel.clockNode := _)
```



# Bringup Flow – JTAG vs TSI



## **JTAG** – standard interface for debug

- Writes program byte-by-byte into target memory using standard JTAG commands
- Send commands to interrupt core, direct core to PC, read core registers, etc.
- Slow, but standard

## **TSI** – tethered serial interface

- Simple fast test chip interface
- Chip communicates with "frontend-server" running on FPGA soft-core



# Chipyard Build System



- Based on Make
- **Variables:**
  - RISCY: specifies directory of riscv-tools installation, should be set by environment file
  - CONFIG: references Scala config describing design
- **RTL-sim Targets:**
  - run-binary-debug: runs BINARY on simulator of CONFIG
- **VLSI Targets:**
  - buildfile/syn/par/drc/lvs: runs steps in VLSI flow
  - redo-syn/par/drc/lvs: runs steps in VLSI flow (does not rerun prereq steps)

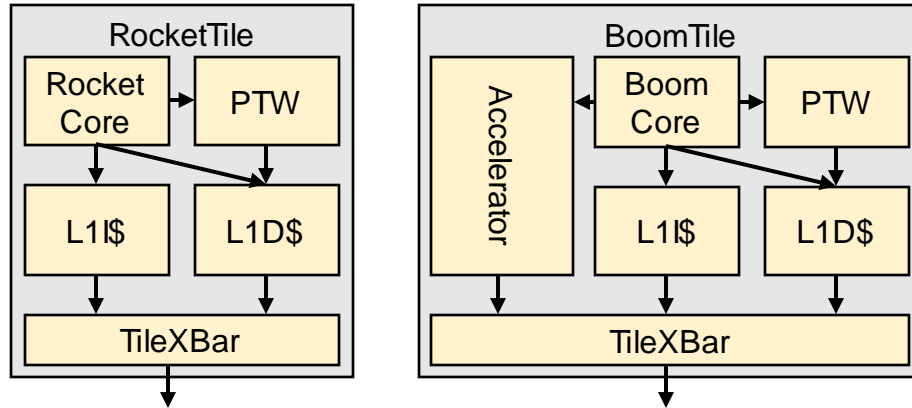


# Chipyard Tiles





# SoC Organization: Tiles



**Tiles:** Units of replication in a multi-core SoC

Contains:


- RISC-V core
- Private L1 caches
- TLBs, PTW
- RoCC accelerator?

Many config options:

- Different core options
- Accelerator interface
- Cache sizes
- Etc.



# ML Tile – Rocket Core



## Rocket:


- First open-source RISC-V CPU
- In-order, single-issue RV64GC core
- Efficient design point for low-power devices
- Virtual memory, floating point, caches, etc.
- Not the focus of this tile, focus on RoCC accelerator interface



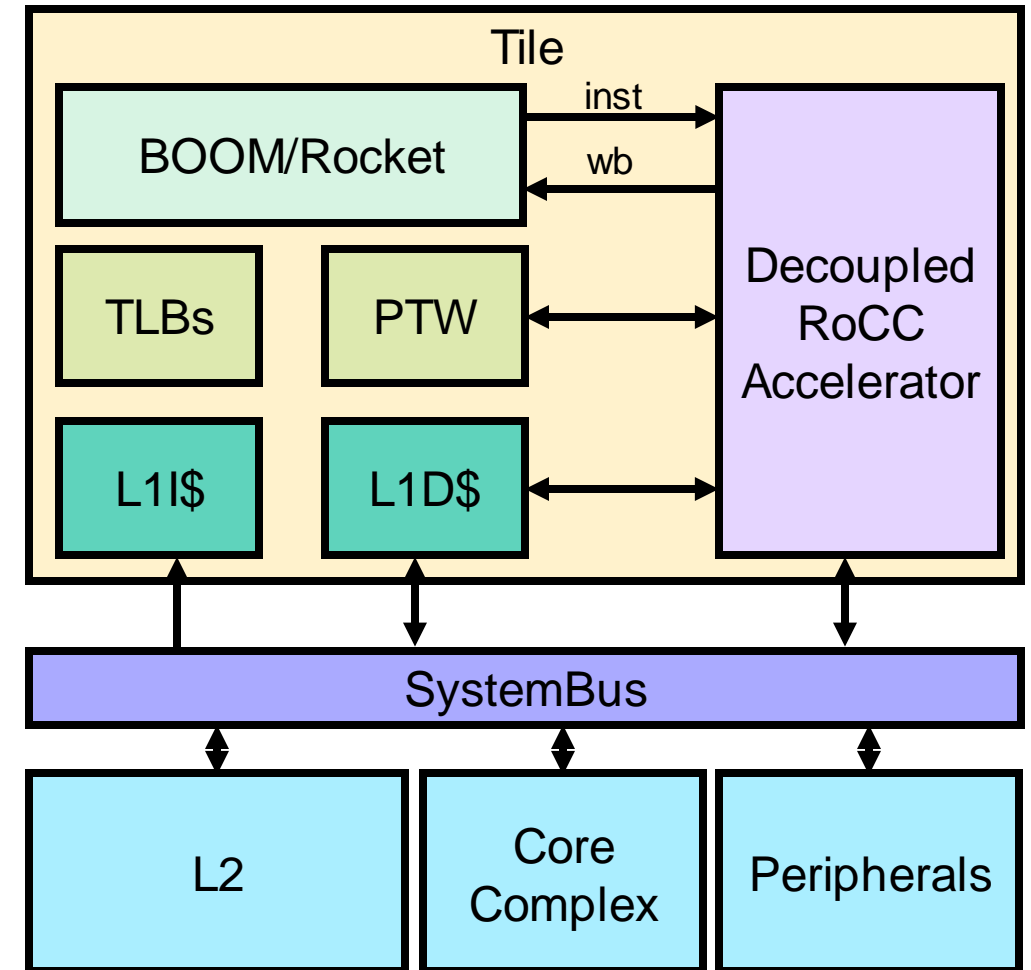
# ML Tile RoCC Interface



- **RoCC:** Rocket Custom Coprocessor
- Sits adjacent to Rocket
- Execute custom RISC-V instructions for a custom extension




- Implement an accelerator as RoCC
  - Receive commands/data from core
  - Access memory through L1D\$
  - Develop SW using RoCC (compare accelerator vs CPU)



# GP Tile – Saturn Core



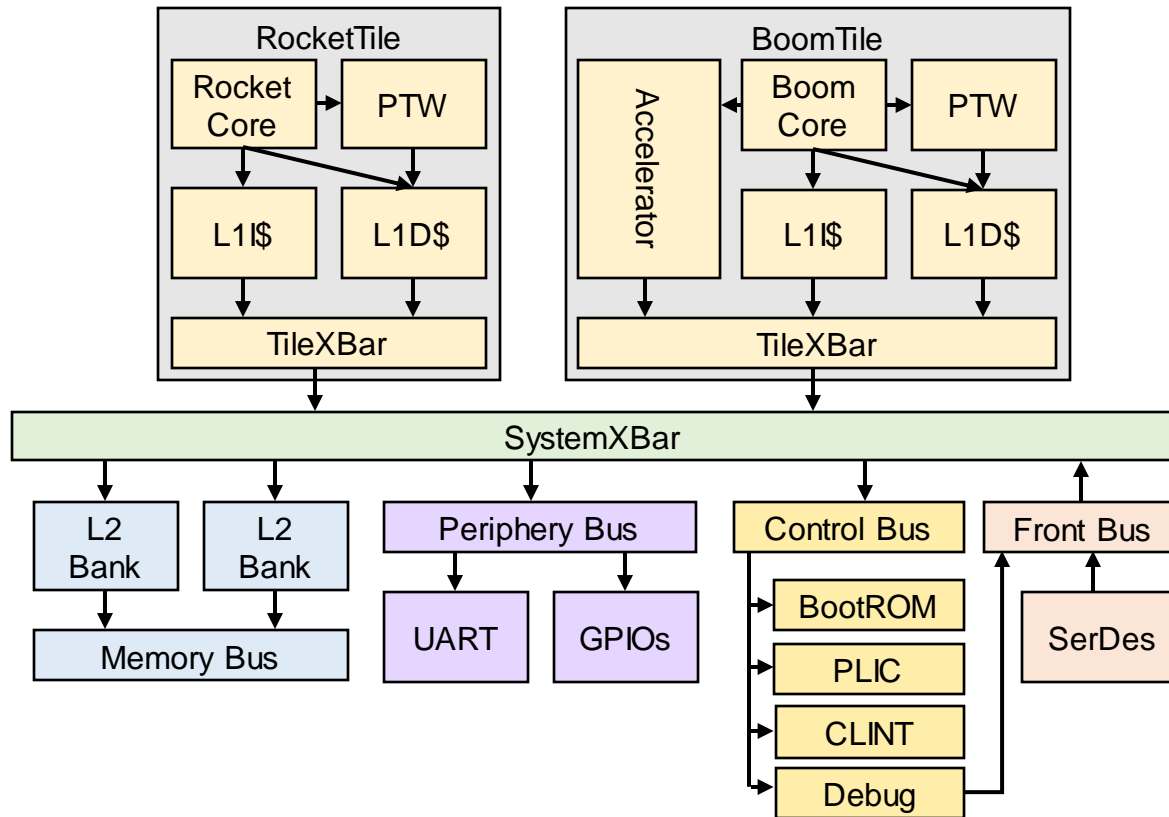
- Superscalar in-order core (fetches/decodes/executes multiple instructions per cycle)
- Baseline: dual-issue (2 instructions per cycle)
- Opportunities for customization
  - Three-issue, four-issue, etc.
  - Advanced cache
  - Custom functional unit
  - Prefetching



# Memory Subsystem



# SoC Organization: Digital System



**L2 Banks:** Distributed banks of L2 cache, each bank manages subset of the address space

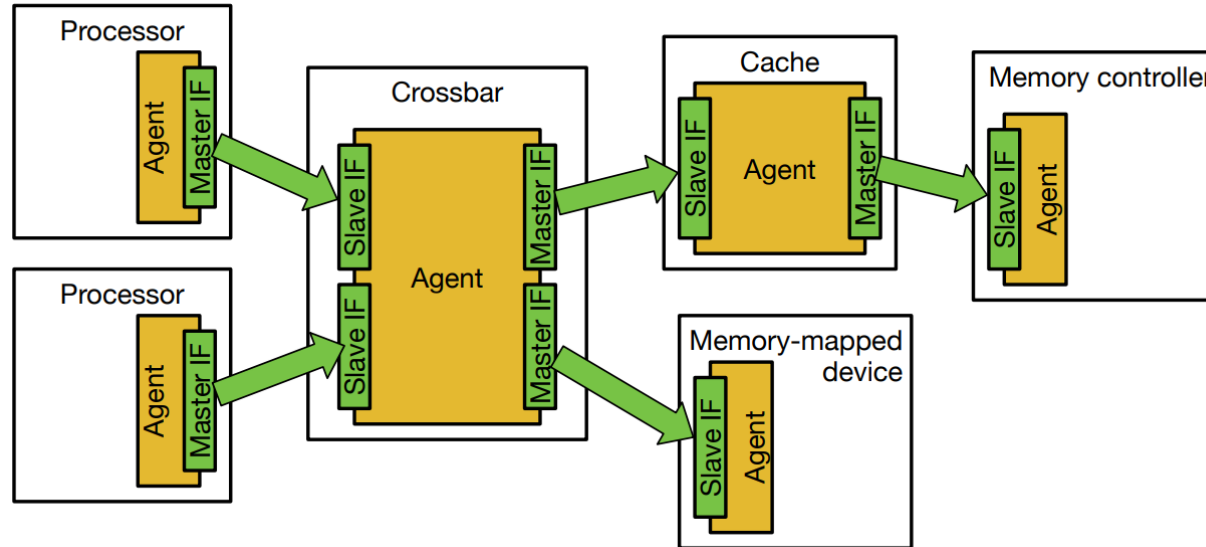
**MMIO devices:** Hang off the system bus or periphery bus

**TileLink:** Coherent interconnect standard (like AXI)

**Interconnect:** Physical transport for interconnect protocol



# TileLink Protocol




- Free and open chip-scale interconnect protocol
- Supports multiprocessors, coprocessors, accelerators, DMA, peripherals, etc.
- Provides a physically addressed, shared-memory system
- Supports cache-coherent shared memory, MOESI-equivalent protocol

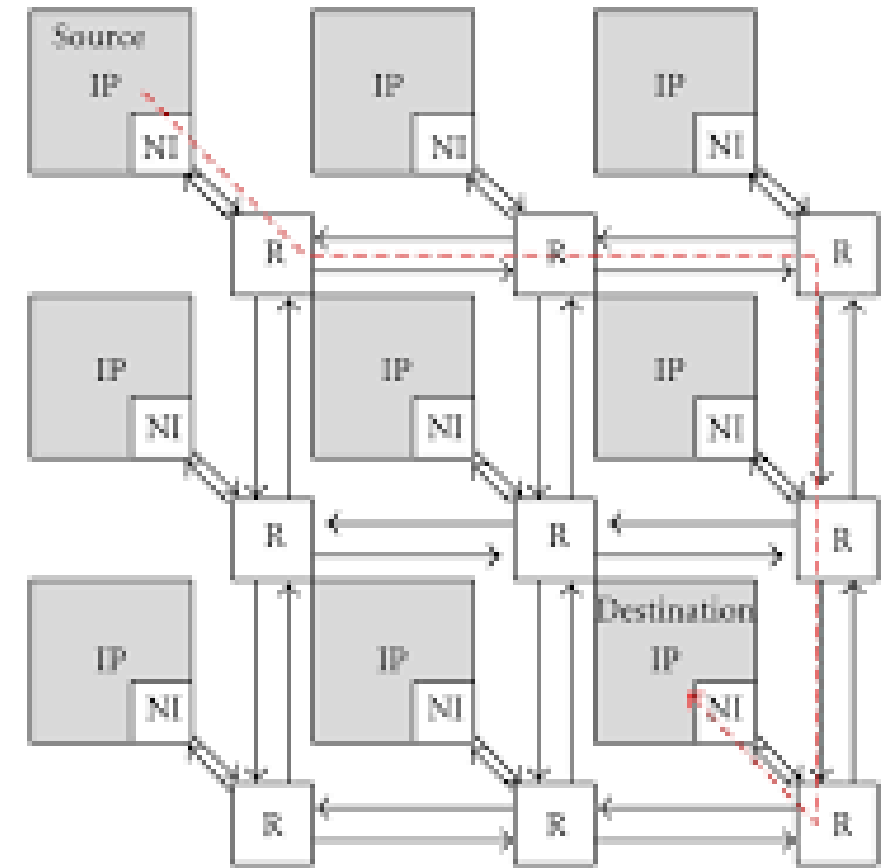




# Coherent Interconnect



- Purpose: manage routing of messages between clients (cores/accelerators/DMA's) and managers (caches/scratchpads/MMIO)
- Baseline: Fully-connected all-to-all crossbar (xbar)
  - Good for small, simple designs
  - Does not scale to large systems
- Goal: Tape-out a network-on-chip
  - Multiplexes message routing onto fewer physical resources
  - We will provide NoC generator



# Questions?

